MLflow in Kubernetes: Streamlined ML Management at Scale

A technical overview of integrating MLflow's experiment tracking and model deployment within Kubernetes environments

 **by Luca Berton**

MLflow Tracking Fundamentals

Log experiments

Record parameters, metrics, and hyperparameters

Store artifacts

Save models, plots, and data samples

Retrieve metadata

Access historical records for analysis and reproduction

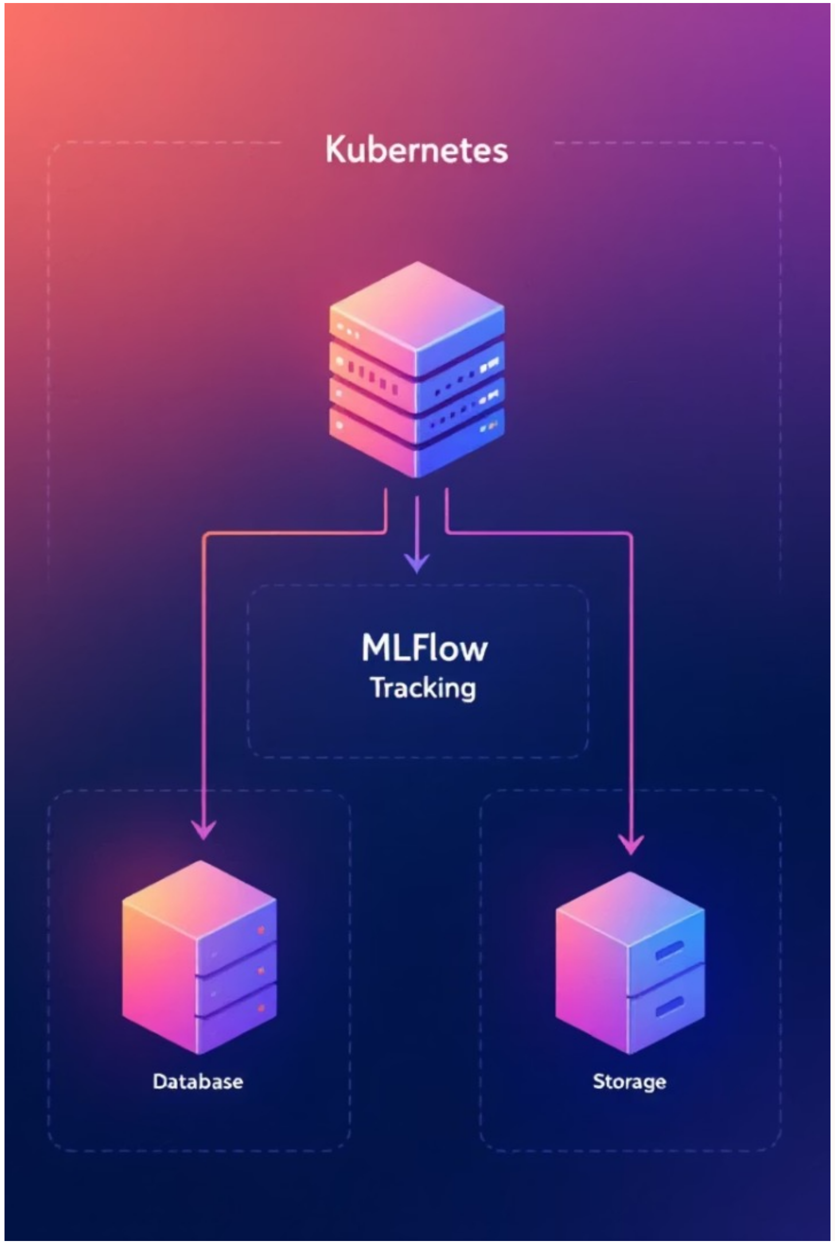

MLflow Tracking in Kubernetes Architecture

Pod Deployment

MLflow server runs in pods managed by Kubernetes Deployment

Persistence Layer

SQL database stores metadata, S3/MinIO/PVC stores artifacts

Service Exposure

Kubernetes Service exposes tracking server internally/externally

Self-Healing

Automatic pod restart if tracking server crashes

Kubernetes Advantages for MLflow Tracking

Horizontal Scaling

Add more pods automatically as experiment load increases

High Availability

Replicated pods ensure tracking remains accessible

Resource Management

CPU/memory limits prevent resource contention

MLflow Models: Structure & Packaging

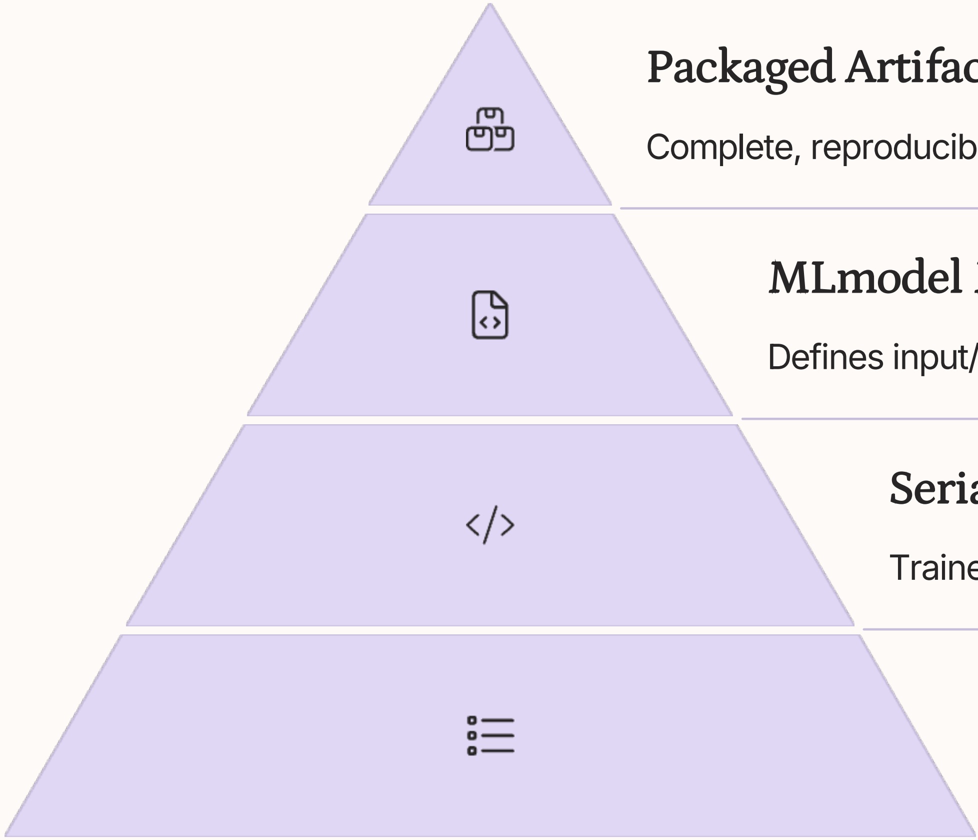

Packaged Artifact

Complete, reproducible model package

MLmodel Metadata

Defines input/output schema and flavor

Serialized Model

Trained model in framework-specific format

Environment Specs

Conda/pip dependencies for exact reproduction

Deploying MLflow Models on Kubernetes

Model Training

Run experiment, track metrics with MLflow

Model Logging

mlflow.log_model() creates portable package

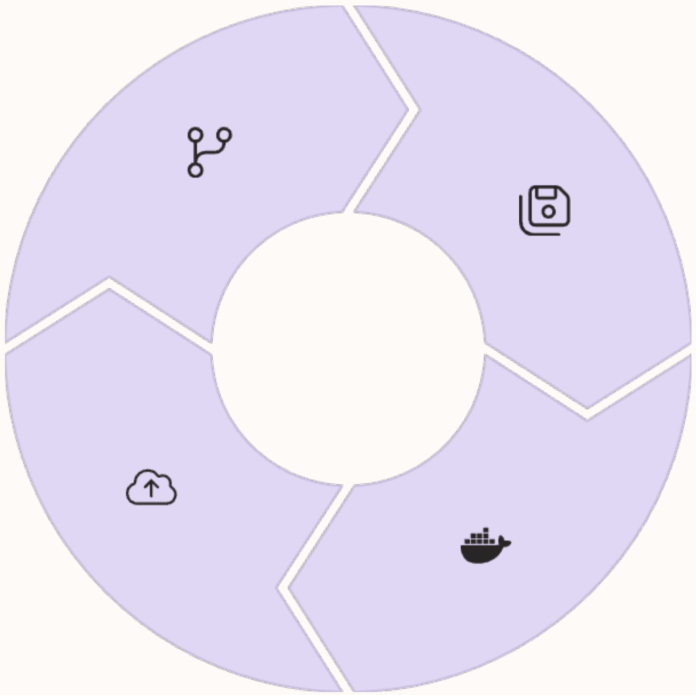

Kubernetes Deployment

Deploy container with KServe or custom resources

Containerization

mlflow models build-docker creates serving image

Model Serving Benefits with Kubernetes

Consistency

- Same model environment everywhere
- Dev-to-production parity
- Reproducible inference behavior

Automation

- Automated rollouts and rollbacks
- Health checks and self-healing
- Automatic load balancing

Observability

- Integrated logging and metrics
- Performance monitoring
- Drift detection integration

Unified ML Lifecycle with MLflow + Kubernetes

Experiment

Track all experiments via MLflow in Kubernetes

Package

Create portable MLflow model artifacts

Deploy

Containerize and serve via Kubernetes

Monitor

Observe performance with integrated tooling